

HTTP

HyperText Transfer Protocol

HTTP

- Created by Tim Berners-Lee at CERN
 - Defined 1989-1991
- Standardized and much expanded by the IETF
- Rides on top of TCP protocol
 - TCP provides: reliable, bi-directional, in-order byte stream



Goal: transfer objects between systems

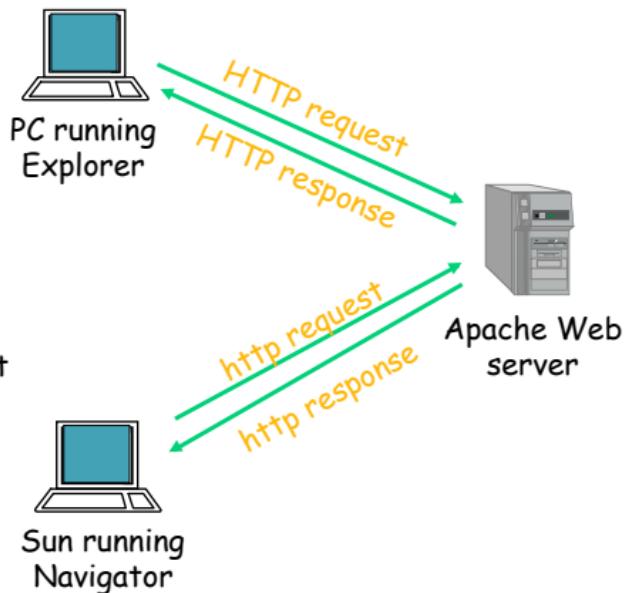
- Do not confuse with other WWW concepts:
 - HTTP is not page layout language (that is HTML)
 - HTTP is not object naming scheme (that is URLs)
- Text-based protocol
 - Human readable

Client-Server Paradigm

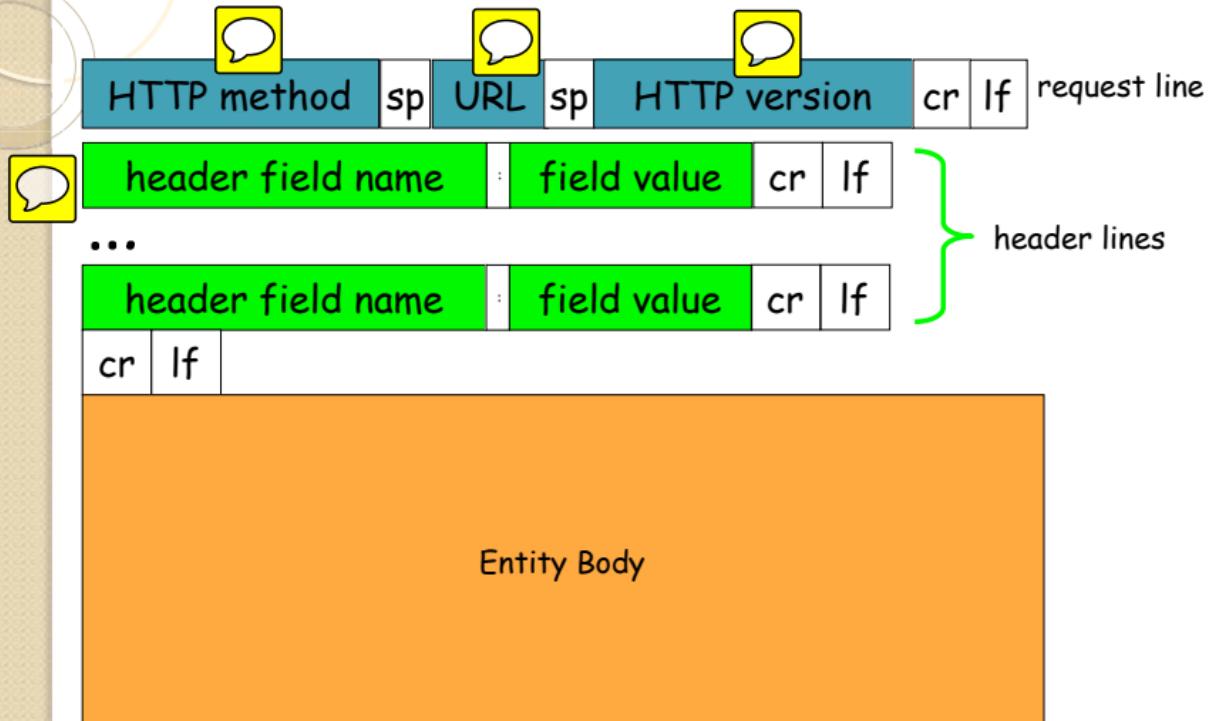


Client:

- ❑ initiates contact with server ("speaks first")
- ❑ typically requests service from server,
- ❑ for Web, client is implemented in browser; for e-mail, in mail reader



HTTP request message



HTTP request format

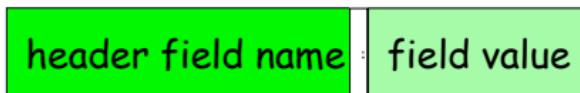
- Request line



- **HTTP method**
 - GET – return content of specified document
 - HEAD – return headers only of GET response
 - POST – execute specified doc with enclosed data
- **URL (only domain portion)**
 - /host-identifier/path
 - e.g. /www.toronto.edu/headlines/
- **HTTP version**
 - e.g. HTTP/1.0

HTTP request format

- Header fields



- Examples:

- Accept: text/html
- Accept: image/jpg
- Accept-language: en; en-gr; fr
- If-modified-since: 17 May 2001
- Content-Length: 2540

HTTP request example



```
GET /somedir/page.html HTTP/1.0
```

```
User-agent: Mozilla/4.0
```

```
Accept: text/html, image/gif, image/jpeg
```

```
Accept-language:fr
```

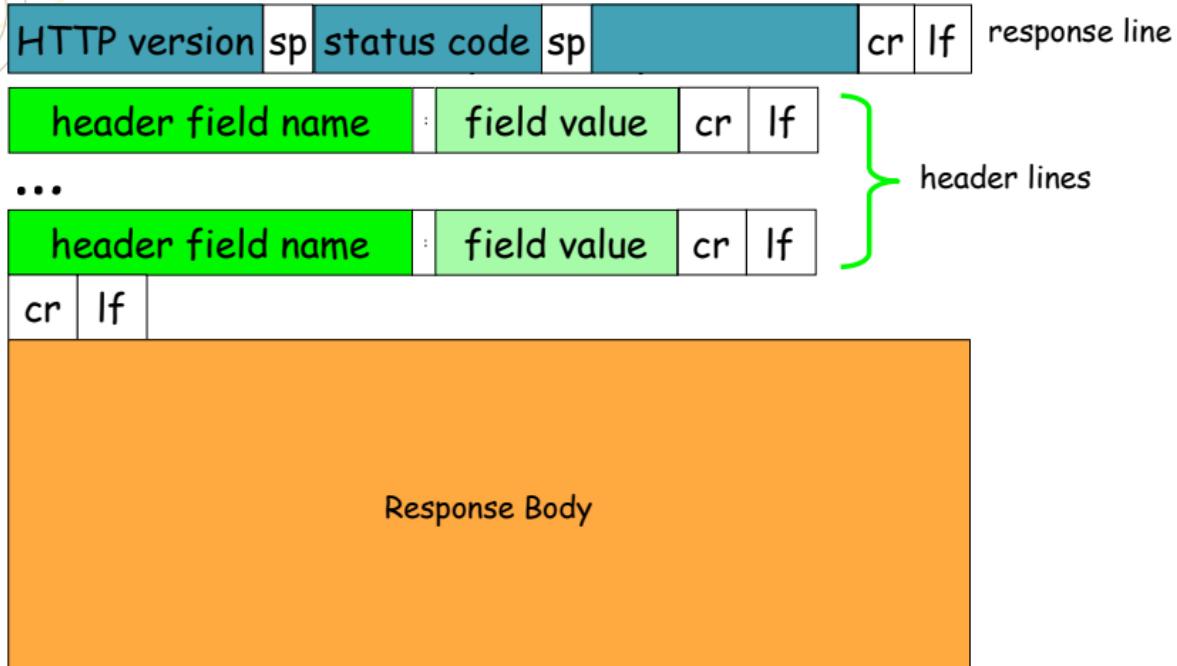
(extra carriage return, line feed)

request line
(GET, POST,
HEAD
commands)

header
lines

Carriage return,
line feed
indicates end
of message

HTTP response message



HTTP response format

- Response line:



- HTTP version
- 3 digit response code
 - IXX – informational
 - 2XX – success
 - 3XX – redirection
 - 4XX – client error
 - 5XX – server error
- Brief text explanation of status code (e.g. OK)

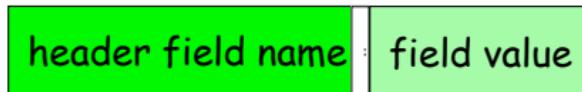
HTTP 1.0: Problems



- Each request opens new connection
- Opening connection takes several packets (why?)
- Starting up is slow (why?)

HTTP response format

- Header fields



- Examples:



- Content-Type: text/html
- Content-Length: 4028
- Language: en;
- Last-modified: 17 May 2004

HTTP response example

HTTP/1.0 200 OK

Date: Thu, 25 Aug 2001 12:00:15 GMT

Server: Apache/1.3.0 (Unix)

Last-Modified: Mon, 22 Aug 2001

Content-Length: 6821

Content-Type: text/html

status line:
(protocol
status code
status phrase)

header
lines

data data data data data ...

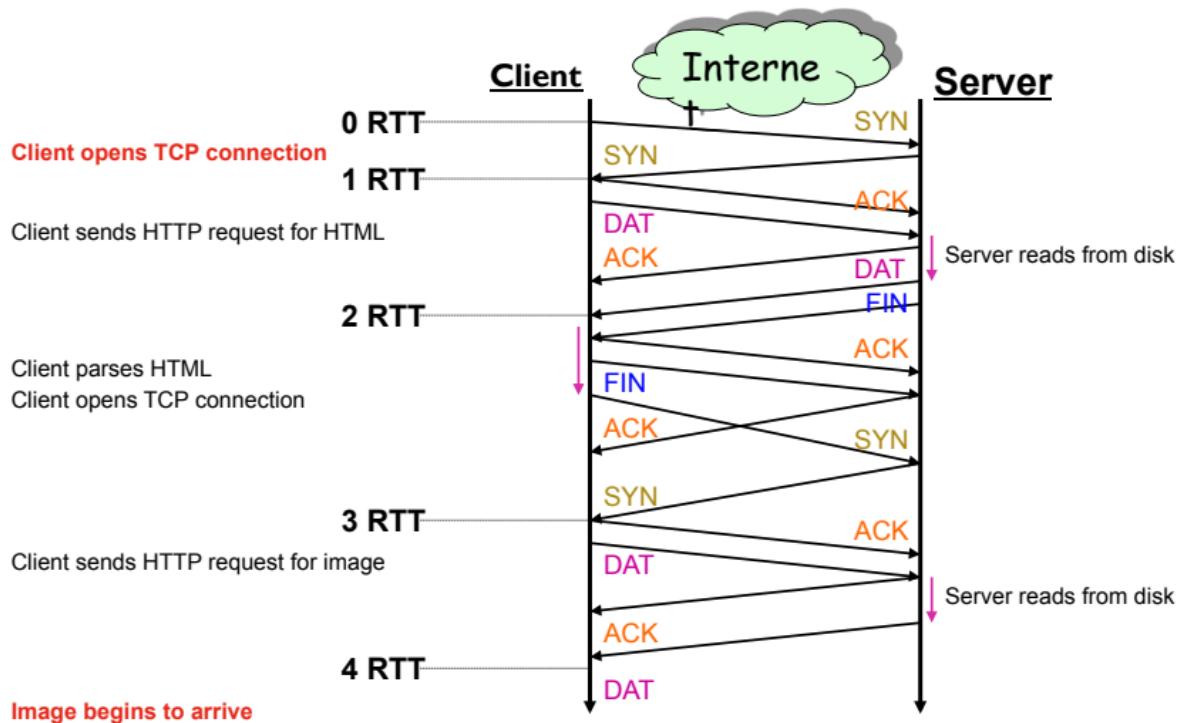
data data data data data ...

data data data data data ...

Carriage return,
line feed
indicates end
of message

data, e.g.,
requested
html file

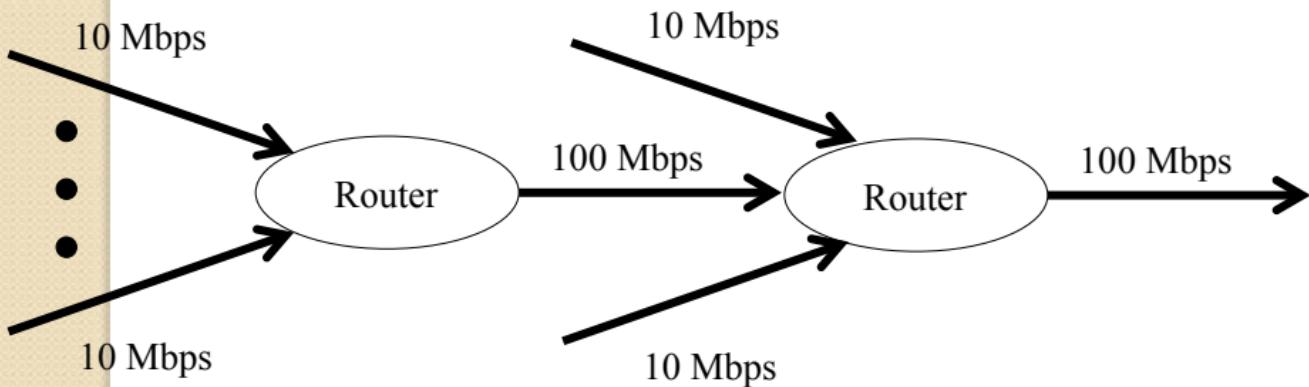
Web Page with Single Image



TCP

- HTTP rides on top of TCP transport service
- TCP provides: reliable, bi-directional, in-order byte stream
- Reliable
 - Prevent packet loss due to congestion
 - Overflow network queues
 - Send at rate at which network can forward packets
 - How to determine sending rate?
 - Dynamic
 - Depends on overall network condition

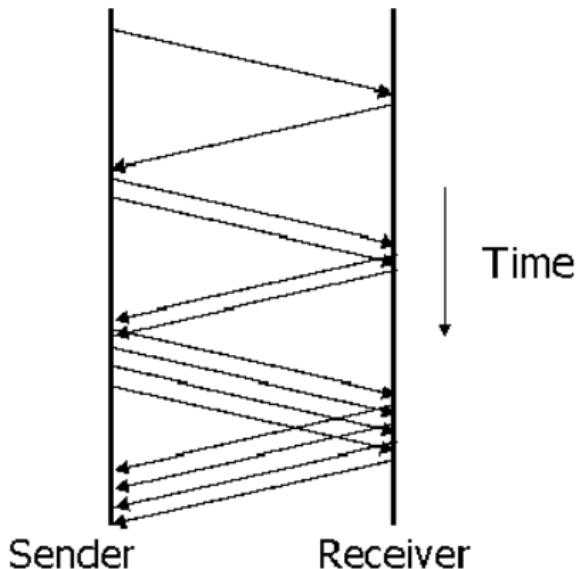
Network Congestion





TCP Slow Start

- Determine sending rate by probing network
- Increase sending rate until a packet is dropped
- Double the number of unacknowledged packets (window size) for every new acknowledgement
- After a drop,
 - Reset window size to 1 packet
 - Cut maximum window size in half
 - Grow window with additive increase



HTTP 1.1: Persistent Connections

- Reuse connection for multiple requests



GET index.html

Connection: keep-alive

... multiple HTTP requests ...

Get banner.gif

Connection: close

Connection length

- When does the data end?
 - Without persistent connections, when connection closes.
 - With persistent connections, reply header includes content length.

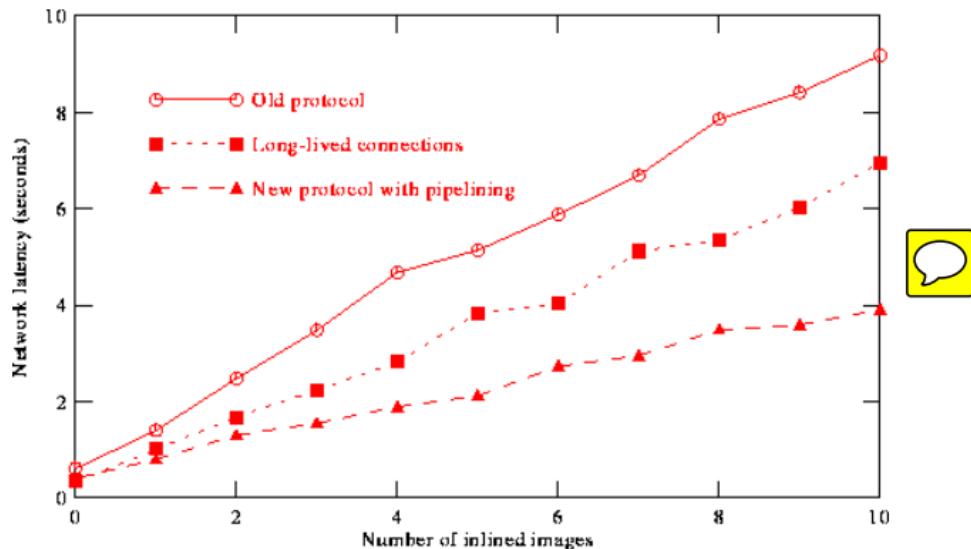
HTTP1.0 vs. HTTP1.1



- Venkata N. Padmanabhan and Jeffrey C. Mogul, "*Improving HTTP Latency.*," in Proceedings of the The 2nd International WWW Conference, Chicago, IL, USA, Oct 1994
- Compared download latency for HTML documents with varying number of embedded images

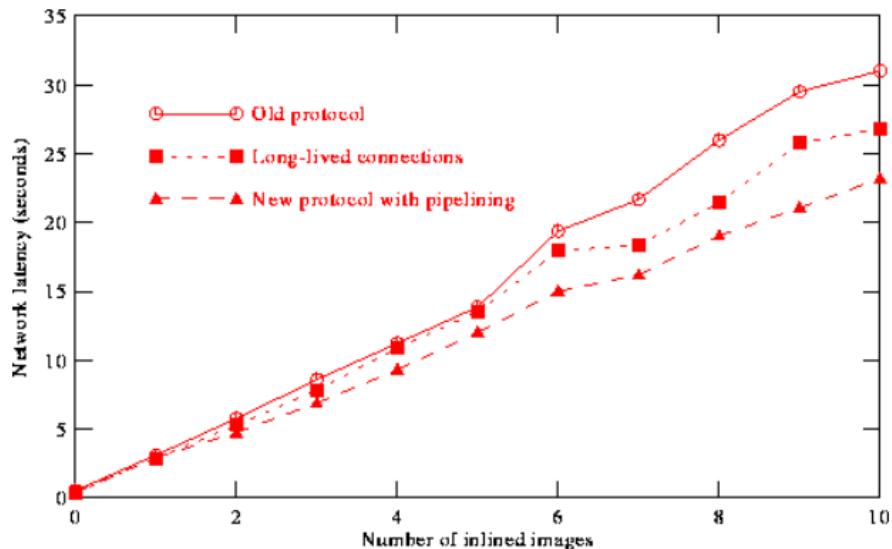
HTTP1.0 vs. HTTP1.1

- Image size 2544 bytes



HTTP1.0 vs. HTTP1.1

- Image size 45566 bytes



Persistent Connection Performance

- Benefits greatest for small objects.
- Serialized requests do not improve response time.
- Pipelining requests can result in large win.
- Server resource utilization reduced due to fewer connection establishments and fewer active connections.
- TCP behavior improved.
 - Longer connections help adaptation to available bandwidth.
 - Larger congestion window improves loss recovery.

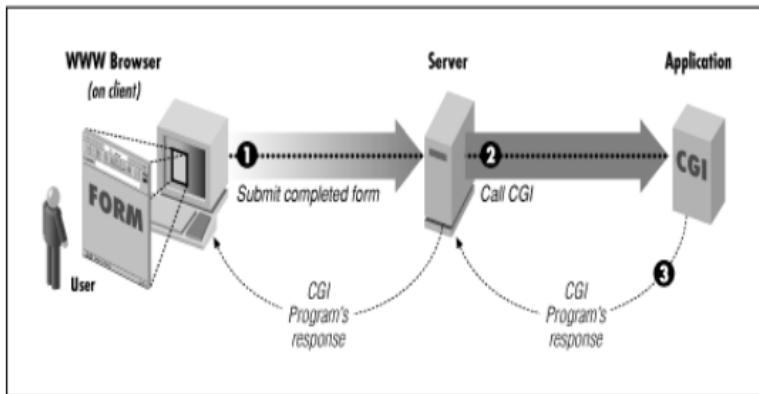


CGI

Common Gateway Interface



CGI Architecture



Objectives

- Standardizes execution of server-side applications/handlers 
- Specifies: 

Web Server-Handler Interaction

- Handler Input
 - Meta-variables 
 - Stdin (message body) 
- Handler Output
 - Stdout 

Example: env.sh

```
#!/bin/bash

echo Content-type: text/plain
echo ""

echo "SERVER_PORT is set to $SERVER_PORT"
echo "The USER_AGENT is $HTTP_USER_AGENT"

echo -----
echo "Values for all environment variables:"""

/usr/bin/printenv

echo -----
```



A screenshot of a web browser window titled 'localhost/cgi-bin/env.sh'. The address bar shows the URL. The browser interface includes tabs for 'CSC 309', 'Most Visited', 'Google Calendar', and 'CSC309 Board'. The main content area displays the output of the env.sh script. It starts with 'SERVER_PORT is set to 80' and 'The USER_AGENT is Mozilla/5.0 (Macintosh; Intel'. Below this, it says 'Values for all environment variables:' followed by a long list of environment variables and their values. The list includes SERVER_SIGNATURE, HTTP_USER_AGENT, SERVER_PORT, HTTP_HOST, DOCUMENT_ROOT, SCRIPT_FILENAME, REQUEST_URI, SCRIPT_NAME, HTTP_CONNECTION, REMOTE_PORT, PATH, PWD, SERVER_ADMIN, HTTP_ACCEPT_LANGUAGE, HTTP_ACCEPT, REMOTE_ADDR, SHLVL, SERVER_NAME, SERVER_SOFTWARE, QUERY_STRING, SERVER_ADDR, GATEWAY_INTERFACE, SERVER_PROTOCOL, HTTP_ACCEPT_ENCODING, REQUEST_METHOD, and HTTP_COOKIE. The output ends with a final '-----'.

```
SERVER_PORT is set to 80
The USER_AGENT is Mozilla/5.0 (Macintosh; Intel
-----
Values for all environment variables:
SERVER_SIGNATURE=<address>Apache/2.2.21 (Unix)
HTTP_USER_AGENT=Mozilla/5.0 (Macintosh; Intel
SERVER_PORT=80
HTTP_HOST=localhost
DOCUMENT_ROOT=/Library/WebServer/Documents
SCRIPT_FILENAME=/Library/WebServer/CGI-Executab
REQUEST_URI=/cgi-bin/env.sh
SCRIPT_NAME=/cgi-bin/env.sh
HTTP_CONNECTION=keep-alive
REMOTE_PORT=52909
PATH=/usr/bin:/bin:/usr/sbin:/sbin
PWD=/Library/WebServer/CGI-Executables
SERVER_ADMIN=you@example.com
HTTP_ACCEPT_LANGUAGE=en-us,en;q=0.7,es-mx;q=0.
HTTP_ACCEPT=text/html,application/xhtml+xml,ap
REMOTE_ADDR::1
SHLVL=1
SERVER_NAME=localhost
SERVER_SOFTWARE=Apache/2.2.21 (Unix) DAV/2 PHP
QUERY_STRING=
SERVER_ADDR::1
GATEWAY_INTERFACE=CGI/1.1
SERVER_PROTOCOL=HTTP/1.1
HTTP_ACCEPT_ENCODING=gzip, deflate
REQUEST_METHOD=GET
HTTP_COOKIE=ConsoleVisibleId0985388-4850-4f39-=
/_=/usr/bin/printenv
-----
```

Client-Handler Interaction



- Clients request the execution of a handler program by means of a:
 - A request method (e.g., GET, POST)
 - Universal Resource Identifier (URI)
 - Identifies handler
 - Extra arguments

URI Syntax



<protocol>://<host><port>/<path-info><script>"?"<query-string>

http://finance.yahoo.com/add?op1=10&op2=20

<protocol>	http
<host>	finance.yahoo.com
<port>	80
<path-info>	null
<script>	add
<query-string>	op1=10, op2=20

Query String Encoding

- RFC 2396
- Why encode? 
 - Can think of a CGI script as a function, send arguments by specifying name/value pairs.
 - Way to pack and unpack multiple arguments into a single string

Encoding Rules

- All arguments
 - Single string of ampersand (&) separated name=value pairs
`name_1=value_1&name_2=value_2&...`
- Spaces
 - Replaced by a plus (+) sign
- Other characters (ie, =, &, +)
 - Replaced by a percent sign (%) followed by the two-digit hexadecimal ASCII equivalent

Method



- GET
 - Arguments appear in the URL after the ?
 - Can be expressed as a URL
 - Limited amount of information can be passed this way
 - URL may have a length restriction on the server
 - Arguments appear on server log
- POST
 - Arguments are sent in the HTTP message body
 - Cannot be expressed as URL
 - Arbitrarily long form data can be communicated (some browsers may have limits (i.e. 7KB)).
 - Arguments usually does not appear in server logs.

Example: add.sh

```
#!/bin/bash

echo "Content-type: text/plain"
echo ""

IFS="&"
set -- $QUERY_STRING

array=($@)

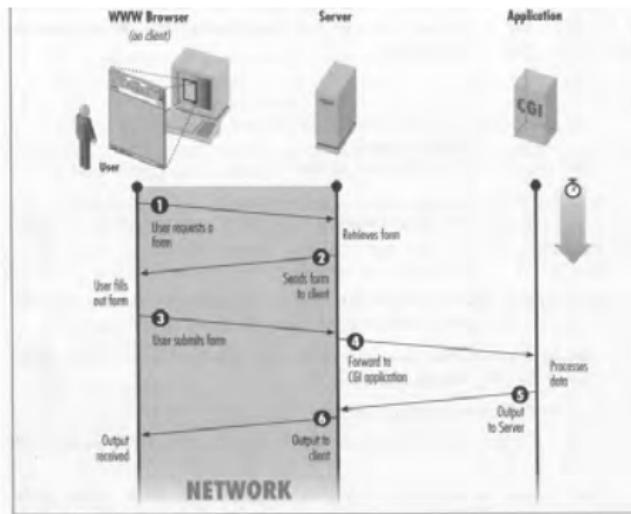
declare -A params

for i in "${array[@]}";
do IFS="=";
set -- $i;
params[$1]="$2";
done

mysum = 0
let mysum+=${params["op1"]}
let mysum+=${params["op2"]}

echo ${params["op1"]} + ${params["op2"]} = $mysum
```

Forms and CGI



- Specify request method in “method” attribute
- Automatically encodes all **named** field values

Example: form.html

```
<!DOCTYPE html>
<html>
  <head>
    <title>form</title>
    <style>
      input { display: block; }
      form { margin: 10px; }
    </style>
  </head>
  <body>
    <h1>GET vs POST</h1>
    <form action="cgi-bin/stdin.sh" method="get">
      Param 1 <input type="text" name="param1" />
      Param 2 <input type="text" name="param2" />
      <input type="submit" value="GET"/>
    </form>

    <form action="cgi-bin/stdin.sh" method="post">
      Param 1 <input type="text" name="param1" />
      Param 2 <input type="text" name="param2" />
      <input type="submit" value="POST"/>
    </form>

  </body>
</html>
```

Example: stdin.sh

```
#!/bin/bash

echo Content-type: text/plain
echo ""

read arguments

echo "STDIN = " $arguments
echo "QUERY_STRING = " $QUERY_STRING
```

Apache CGI

- Apache directories
 - cgi-bin cgi files
 - htdocs .html, .gif, .jpeg, .css, .js
 - conf configuration files
- httpd.conf
 - Main Apache configuration file
 - CGI are disabled by default
 - To enable CGI:
 - Add "ExecCGI" to the "Options" directive

Options Indexes FollowSymLinks **ExecCGI**